Assignment Project Exam Help

https://powcoder.com

Add WeChat powcoder



EEEE4115 Advanced Assignment Protect Exam Help Computational https://gineeringm

Introduction to GPU Programming with CUDA®

Add WeChat powcoder



CUDA Programming Model

 The GPU is seen as a 'compute device' to execute all or part of an Assignment Project Exam Help application that:

- Can be isolated as a function
- Has to be executed med Windlest powcoder
- Is highly data parallel
- Is lightweight

• The function when compiled to run on the 'device' is called a 'Kernel'



CUDA Programming Model

- Allocate memory on device
 - Allocated by cuda Mansignment Project Exam Help
- Copy input arrays from host to device memory bowcoder.com
 - cudaMemcpy()
- Execute kernel(s) on the device WeChat powcoder
 - kernel <<< block size, number of blocks, shared memory size
 >>(parameters);
- Copy output arrays from device to host
 - cudaMemcpy()



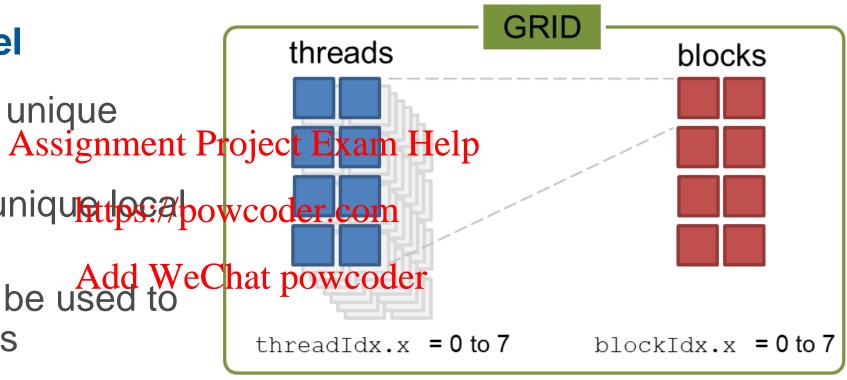


Thread/Block Model

Each thread has a unique local index.

■ Each block has a uniquettocapowcoder.com index.

 These indices can be used to Chat powcoder index data in arrays



kernel <<<8,8>>> (..);

Thread/Block Model

```
int data [10] = { 2, 4, 6, 8, 10, 12, 14, 16, 18, 20}
                         Assignment Project Exam
creates:
                              https://powcoder.com
to select i<sup>th</sup> element:
                     int i = Add We Chat powcoder
kernel <<<1,10>>>:
                      do something with data[i];
                    int i = blockIdx.x;
kernel <<<10,1>>>:
                      do something with data[i];
                     int i = blockDim.x * blockIdx.x + threadIdx.x;
kernel <<<2,5>>>:
                      do something with data[i];
(kernel <<<number of blocks, number of threads per block>>>(...);
```



```
a[] = \{1, 2, 3, 4, 5, 6, 7, 8, 9, ...\} (number of elements = int length)
       a:
             2 3
                     5
                        6 | 7 |
                         Assignment Project Exam Help
CPU code to access elements of a []:
                              https://powcoder.com
          void foo (int *a, int *b) {
              for (int i=0; iAddtWeChat powcoder
              b[i] = a[i];
```



```
a[] = {1,2,3,4,5,6,7,8,9,...} (number of elements = int length)

a: 1 2 3 4 5 6 7 8 9
```

GPU code to access elements of signing entite Representation of the code to access elements o

```
https://powcoder.com

a: 1 2 3 4 5 6 7 8 9 WeChat powcoder

thread no. 1 2 3 4 5 6 7 8 9

__global___ void foo(int *a, int *b) {
    int index = threadIdx.x; //unique thread ID
    if (index < length)
        b[index] = a[index];
}
```

Remember, max. number of threads = 512 or 1024 (hardware dependant)



```
= \{1, 2, 3, 4, 5, 6, 7, 8, 9, ...\} (number of elements = int length)
GPU code to access elements of signing entities of single transfer of signing entities of signing entities
                               block no.
                                                                                                                                                                                                                                            https://powcoder.com
                                                                                                                                                                                                                                       Add WeChat powcoder
                        thread no.
                                                               global void foo(int *a, int *b) {
                                                                                     int index = blockIdx.x; //unique block ID
                                                                                     if (index < length)</pre>
                                                                                                                   b[index] = a[index];
```

Remember, max. number of blocks = 65535



```
a[] = \{1, 2, 3, 4, 5, 6, 7, 8, 9, ...\} (number of elements = int length)
GPU code to access elements of a justing blocks of n threads. Help
e.g. n = 3:
                                  https://powcoder.com
             block no 1
                        block no 2
                                  Add WeChat powcoder
        a:
    thread no.
                          2 3
         global void foo(int *a, int *b) {
            int index = threadIdx.x+blockIdx.x*blockDim.x;
            if (index < length)</pre>
                b[index] = a[index];
```



Assignment Project Exam Help

https://p&xalmpla: Computing T

Add WeChat powcoder

Illustrating memory types



Computing π

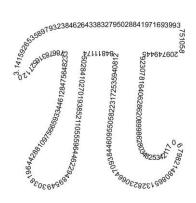
Approach

3 tier reduction:

- Every thread adds Assignment Project Exam Help $\int_0^1 \frac{1}{1+x^2} \approx 4\Delta x \sum_{i=0}^1 \frac{1}{1+x_i^2}$
- The first thread in each block adds all results from the block from the block
- The CPU does the remaining weder tions wooder

3 types of memory

- Each thread uses private memory to store their partial results
- The first thread of a block needs access to all results in the block using shared memory
- The block totals need to by copied to CPU global memory





Computing π - Device Code

```
global void pi(float *blockSums, int stepsPerThread, float dx) {
12
13
       shared float threadSums[BLOCKSIZE];
14
15
       int id
              = threadIdx.x + blockDim.x * blockIdx.x;
                                                           allocate work to each thread
       int istart = id * stepsPerThread;
16
       int istop = istart + step Assignment Project Exam Help
17
18
19
       float accum = 0.0f;
20
       for (int i = istart; i < istop; i https://powcodeream thread calculates their
22
       threadSums[threadIdx.x] = accum; Add WeChat powcoder contribution to overall approximation
23
24
25
26
       __syncthreads();
27
                                                            wait for all threads to finish
28
29
       if (threadIdx.x == 0) {
30
           float blockSum = 0.0f;
                                                            calculate sum of contributions from
31
           for (int j = 0; j < blockDim.x; <math>j++) {
32
              blockSum += threadSums[j];
                                                            threads in each block
33
34
           blockSums[blockIdx.x] = blockSum;
35
36
```



Computing π - Host Code

```
39 int main()
40 □{
41
        cudaError t err;
42
43
        const int stepsPerThread = 2048;
44
        const int blockSize = BLOCKSIZE;
45
        const int numBlocks = 256;
46
47
        const int numSteps = blockSize * numBlocks * stepsPerThread;
48
        const float dx = 1.0f / numSteps;
                                             ssignment Project Exam Help
49
        float *d blockSums;
51
        err = cudaMalloc((void**)&d_blockSums, sizehftips://powcoder.com
assert(err == cudaSuccess);
53
54
55
        err = cudaMemcpy(d blockSums, h blockSums, sizeof(float) * numBlocks, cudaMemcpyHostToDevice);
56
57
            assert(err == cudaSuccess);
                                                       d WeChat powcoder
58
        pi<<<numBlocks, blockSize>>> (d blockSums, stepsPerThread, d
59
60
61
        err = cudaMemcpy(h blockSums, d blockSums, sizeof(float) * numBlocks, cudaMemcpyDeviceToHost);
62
            assert(err == cudaSuccess);
63
64
        float pi = 0.0f;
65
        for (int i = 0; i < numBlocks; i++)
66
            pi += h blockSums[i];
67
68
        pi *= dx;
69
        printf("pi approximately equals: %f\n", pi);
71
72
        cudaFree(d blockSums);
73
        free(h blockSums);
74
        return 0;
76 4
```

assign total amount of work per thread/block

allocate memory on device, copy arrays and check for errors call kernel copy sums from each block sum all



Assignment Project Exam Help

Example: Computational Electromagnetics

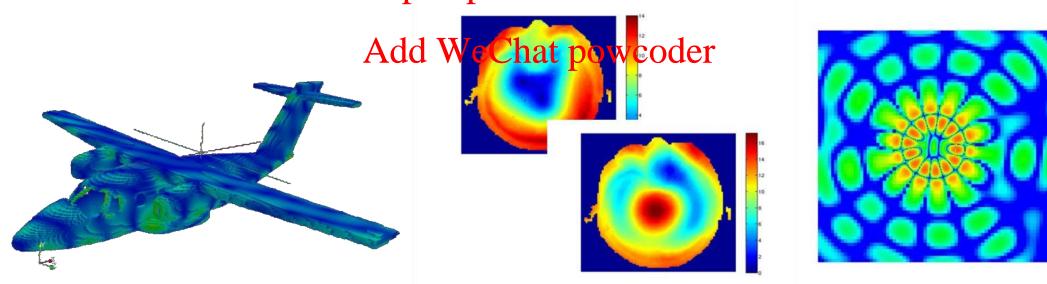
Add WeChat powcoder

Porting the 1D TLM Algorithm to CUDA



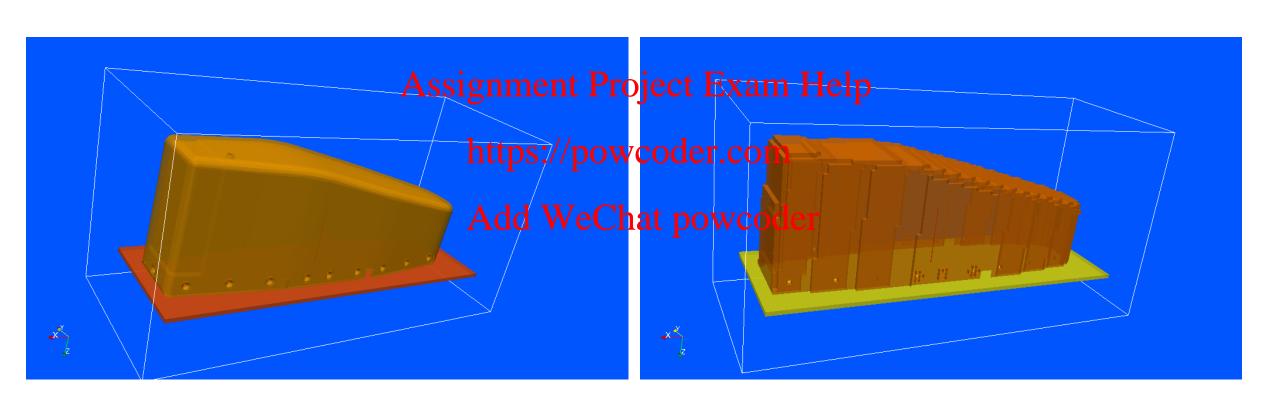
Transmission Line Method

- Founded and developed within the University for over 40 years
- Time domain method for the computation of electromagnetic fields
- Problems are discretised in time and space
- The computational domain being broken down into a mesh of transmission line segments connected at 'nodes'
 Assignment Project Exam Help
- Extremely powerful and underpins the simulation strategies employed commercial software https://powcoder.com





TLM Example: Aircraft Radome





TLM Example: Aircraft Radome





The Algorithm - Consider a line in free space i.e. a wire.

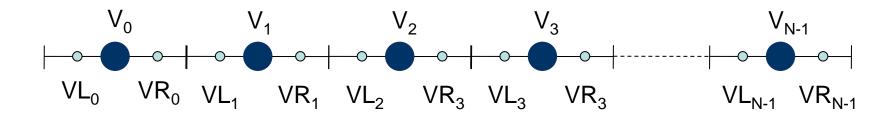
Problem space subdivided into N equal sections (nodes) . Simulation is then run for a total time T, over a number of time steps of Δt .

Assignment Project Exam Help

A simplified description of the process. At each Δt , two operations occur in sequence:

- 1. VL_n and VR_n are obtained from https://powcoder.com
- 2. VR_n and VL_{n+1} and are swapped The 'connect' process der
- 3. The new V_n values are calculated

Operations are highly data parallel!





1st Things 1st

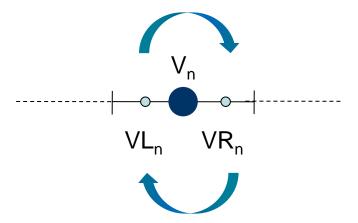
- Establish Benchmarks:
 - Execution Time for CPU application
 - Output from CPUAspignment Project Exam Help

https://powcoder.com

Add WeChat powcoder











```
//allocate memory on device
cudaMalloc((void**)&dev VR, N*sizeof(float));
cudaMalloc((void**)&dev VL, N*sizeof(float));
cudaMalloc((void**)&dev Vt, NT*sizeof(float));
//copy memory areas Assignment Project Exam Help
NT*sizeof(float), cudaMemcphttps://powcoder.com
for (int m = 0; m < NT; m+Add WeChat powcoder
```

Allocate memory on device



```
//copy memory areas from host to device cudaMemory (dev_VR, VASSIZEMMENT, Purple Cto Exame He Copy memory areas
cudaMemcpy(dev VL, VL, N*sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(dev Vt, Vt, NT*sizeof(float), cudaMemcpyHostToDevice);
    float source = tlmSourcAdd, WeChat) powcoder
    tlmScatter <<<1, N>>> ( dev VR, dev VL, N, m, source);
```

from host to device



```
cudaMemcpy (dev_VR, Assignment Project Exam Helpaunched as two
N*sizeof(float), cudaMemcpyhttps://powcoder.com
//run TLM algorithm for NTAdd WeChat powcoder
for (int m = 0; m < NT; m++)
   float source = tlmSource(m*dt, delay, width);
   tlmScatter <<<1,N>>>( dev VR, dev VL, N, m, source);
   tlmConnect <<<1,N>>> ( dev VR, dev VL, dev Vt, N, m);
```

Launch kernels on device

Note: for generality scatter & connect are separate kernels to ensure the scatter completes before the connect.

If using single block of threads then both routines could be in the same kernel



```
//allocate memory on device
                                                                 global void tlmScatter (float* VR,
                                                                float* VL, int N, int m, float source)
cudaMalloc((void**)&dev VR, N*sizeof(float));
cudaMalloc((void**)&dev VL, N*sizeof(float));
                                                                     unsigned int idx = threadIdx.x;
                                                                     //apply source
cudaMalloc((void**)&dev Vt, NT*sizeof(float));
                                 Assignment Project Exam Helpce;
//zero arrays .....
                                        https://powcoder.fcom N)
//copy memory areas from host to device
                                                                         float V = VL[idx] + VR[idx];
cudaMemcpy(dev VR, VR, N*sizeof(float),cudaMemcpyHos
                                                                         VR[idx] = V - VR[idx];
cudaMemcpy(dev_VL, VL, N*sizeof(float),cudaMemcp
cudaMemcpy(dev_Vt, Vt, NT*sizeof(float),cudaMemcpyHostToDevice);
//run TLM algorithm for NT times steps
for (int m = 0; m < NT; m++)
  float source = tlmSource(m*dt, delay, width);
                                                                * Kernel calls for each timestep m are queued
  tlmScatter <<<1,N>>>( dev VR, dev VL, N, m, source);
  tlmConnect<<<1,N>>>( dev VR, dev VL, dev Vt, N, m);
                                                                on device
```

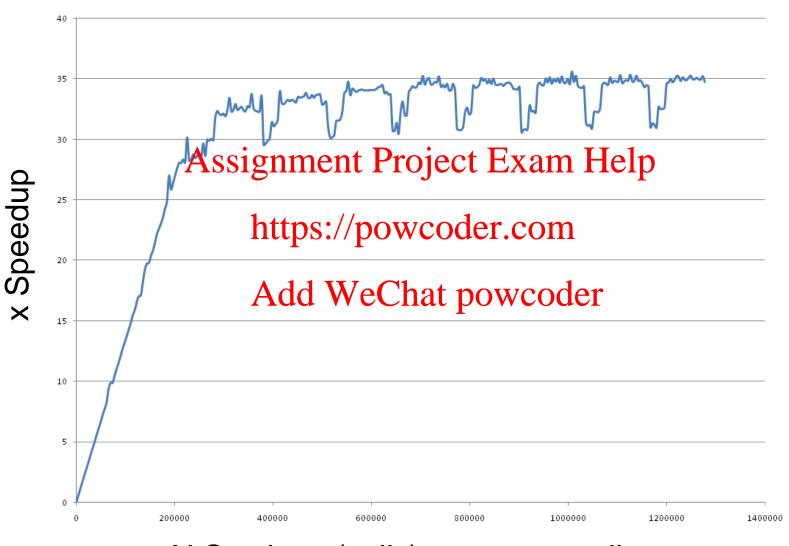


```
//allocate memory on device
                                                                                                                                                                                                                global void tlmConnect (float* VR,
                                                                                                                                                                                                           float* VL, float* Vt, int N, int m)
cudaMalloc((void**)&dev VR, N*sizeof(float));
cudaMalloc((void**)&dev VL, N*sizeof(float));
                                                                                                                                                                                                                        unsigned int idx = threadIdx.x;
cudaMalloc((void**)&dev_Vt, NT*sizeof(float));
                                                                                                                                                                                                                        //Connect
                                                                                                    Assignment Project Exam Helpidx < N)
//zero arrays .....
                                                                                                                                                                                                                                      float V = VR[idx-1];
                                                                                                                                                                                                                                     VR[idx-1] = VL[idx];
                                                                                                                           https://powcoder.comidx] = V;
//copy memory areas from host to device
cudaMemcpy(dev_VR, VR, N*sizeof(float),cudaMemcpyHostToDevice);
\verb|cudaMemcpy| (dev_VL, VL, N*size of (float), \verb|cudaMerArd - | cut | vere | hat | pow code | leaves 
cudaMemcpy(dev_Vt, Vt, NT*sizeof(float),cudaMemcpyHostToDevice);
                                                                                                                                                                                                                        if (idx == 0)
                                                                                                                                                                                                                                     VR[N-1] *= -1.f;
//run TLM algorithm for NT times steps
                                                                                                                                                                                                                                     VL[0] = 0.f;
for (int m = 0; m < NT; m++)
                                                                                                                                                                                                                                     Vt[m] = VL[2] + VR[2]; // o/p
      float source = tlmSource(m*dt, delay, width);
      tlmScatter <<<1,N>>>( dev VR, dev VL, N, m, source);
      tlmConnect<<<1,N>>>>( dev VR, dev VL, dev Vt, N, m);
```



```
global void tlmScatter (float* VR, float* VL, int N, int m, float
source) {
   unsigned int idx = threadIdx.x + blockIdx.x*blockDim.x;
   //apply source
   if (idx == 0) //only one thread to apply source
       VL[0] += source;
   //scatter
       (idx < N) {
float V = VL[idx] + VR[idx] Assignment Project Exam Help
   if (idx < N) {
       VR[idx] = V - VR[idx];
       VL[idx] = V - VL[idx];
                                       https://powcoder.com
                                       Add WeChat powcoder
int main () {
unsigned int nBlocks = (N+BLOCKSIZE-1)/BLOCKSIZE;//Ensure sufficient number of blocks
   for (int m = 0; m < NT; m++)
        float source = tlmSource(m*dt, delay, width);
        tlmScatter<<<nBlocks, BLOCKSIZE>>> ( dev VR, dev VL, N, m, source);
        tlmConnect<<<nBlocks,BLOCKSIZE>>> ( dev VR, dev VL, dev Vt, N, m);
```





N Sections (cells) to represent line



Assignment Project Exam Help

https://powcodettppA Coursework 2

Add WeChat powcoder

Porting the 1D TLM Algorithm to CUDA



CUDA Coursework 2

Coursework is worth 25% of the Module

- Full submission to include the full application code as a single .cu file and a brief report (approximatelynte) pages that and the full submission to include the full application code as a single .cu file and
 - A benchmark of the provided 2D TLM code on your platform of choice (illustrative result and rulations) powcoder.com
 - A discussion that identifies the areas of the code that are data parallel
 A description of the scatter and connect functions as implemented on the
 - A description of the scatter and connect functions as implemented on the GPU Obtain performance results
 - A discussion of any routes taken to optimise the application
 - A benchmark of the GPU application (illustrative result, runtime and speed up compared to CPU version)
- Coursework deadline 11th January 2021 at 15:00